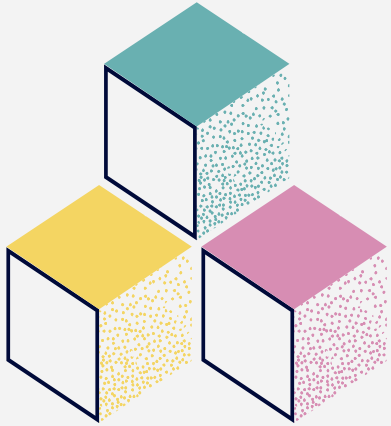

Predicting Kickstarter Project Success

Erick Cantu Perez, Su Wong
30.06.2022

Our stakeholder

Mr. Jacoby Stevens

- Retired entrepreneur
- Likes supporting crowdfunding projects
- Wants to give back to the community
- Wants to support the projects most likely to succeed
- Would like us to predict for him the likelihood of success or failure

About Kickstarter

- Founded in 2009, is a crowdfunding website
- Has an all-or-nothing funding model
- A project is only funded if it meets its goal amount; otherwise no money is given by backers to a project

About the dataset

- Dataset consists of 209222 projects between April 22 2009 and March 14 2019.

- 5 different outcomes for projects on Kickstarter: successful, failed, live, canceled, suspended. Only successful and failed projects were used to train our model.

- Taking only successful/failed projects and eliminating duplicate projects left us with 168979 projects.

Exploratory Data Analysis findings

The predictors for our model are:

- Category (e.g. Art, Food, Technology)
- Country
- Fundraising goal in USD
- Campaign Duration
- Was the project selected by the staff as “staff pick”

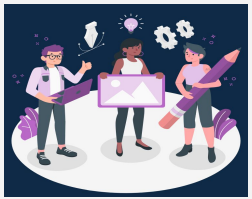

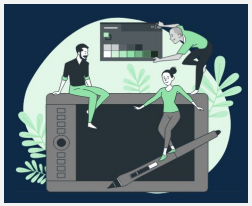

Exploratory Data Analysis findings

Overall success rate of all projects is 56%

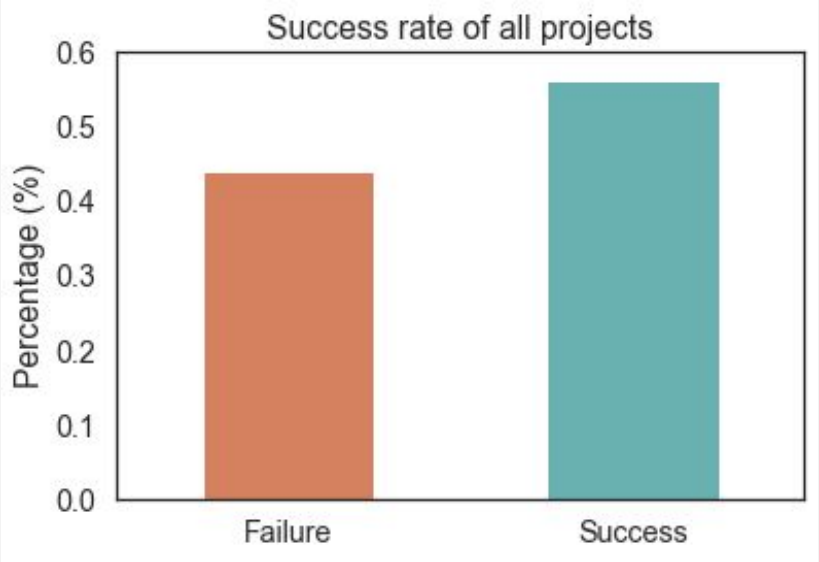

Goal ranged between 0.1 and 15.23M USD($)

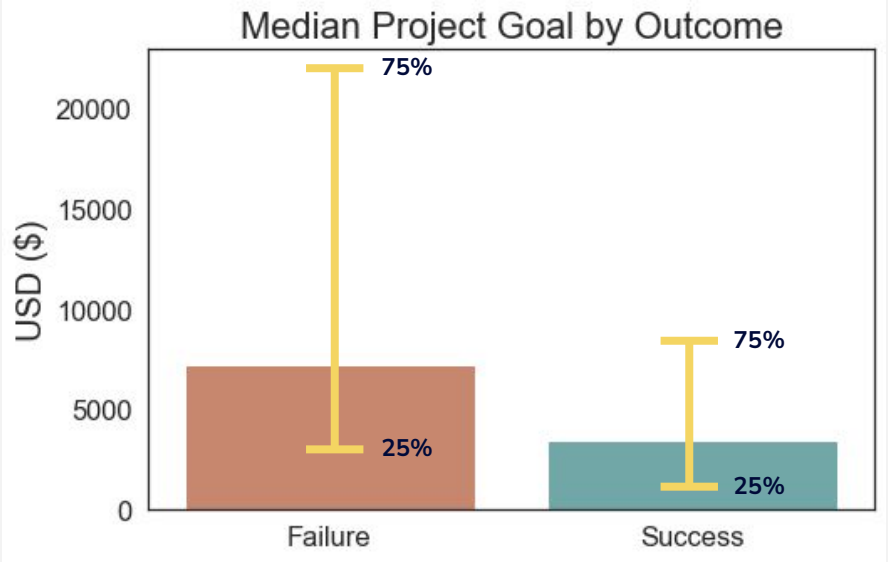

Baseline Model

Baseline Model:

We think that projects based in USA are likely to be successful

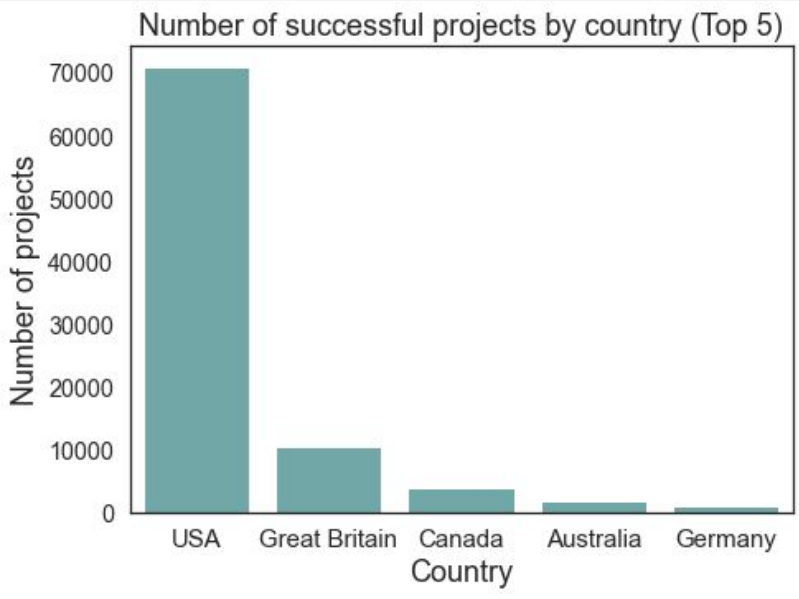

Evaluation Metric:

We train the model based on precision. Precision is the ratio of *correctly* predicted positives out of *all* of the results that were predicted positive.

Score for baseline model:

Precision = 57.8%, Accuracy* = 55.2%

*Accuracy is the percentage of all correctly predicted outcomes

Machine Learning models

Because our data has more categorical variables, we decided to focus more on tree-based supervised machine learning models:

1. Logistic regression
2. Decision Trees
3. Random forests
4. XGBoost
5. Adaboost

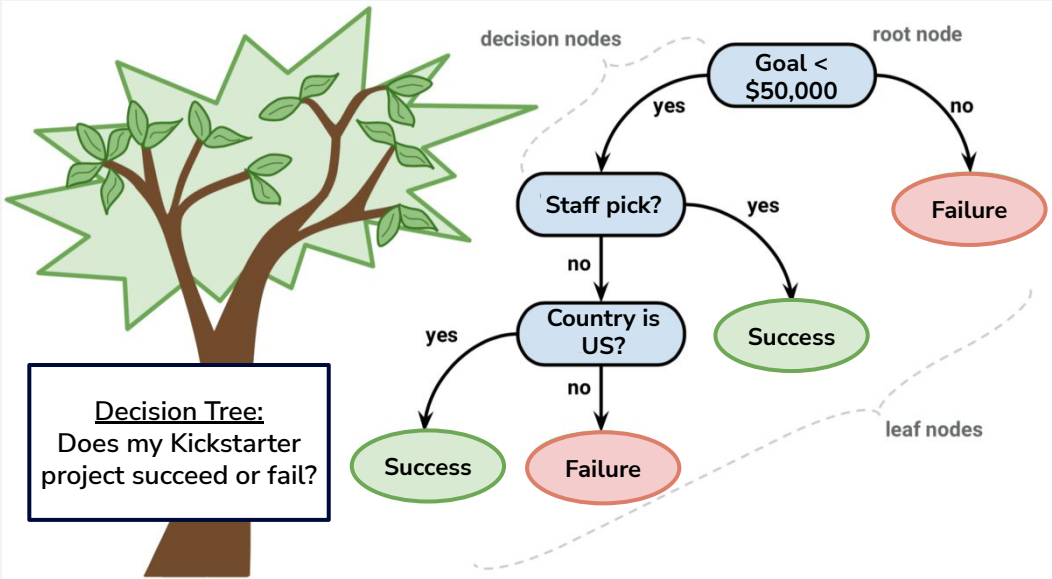

Results

XGBoost gives the best results “out of the box”

	Logistic Regression	Decision Tree	Random Forest	XGBoost	Adaboost
Precision*	69.5%	70.0%	70.5%	72.2%	69.4%
Accuracy**	69.6%	66.6%	68.4%	72.1%	67.7%

After optimizing the parameters:
Final precision: **72.1%** (Accuracy: **71.9%**)

* Precision is the ratio of *correctly* predicted positives out of *all* of the results that were predicted positive.
** Accuracy is the percentage of all correctly predicted outcomes

Error Analysis

Our model is only 72% accurate. How can we improve this number?

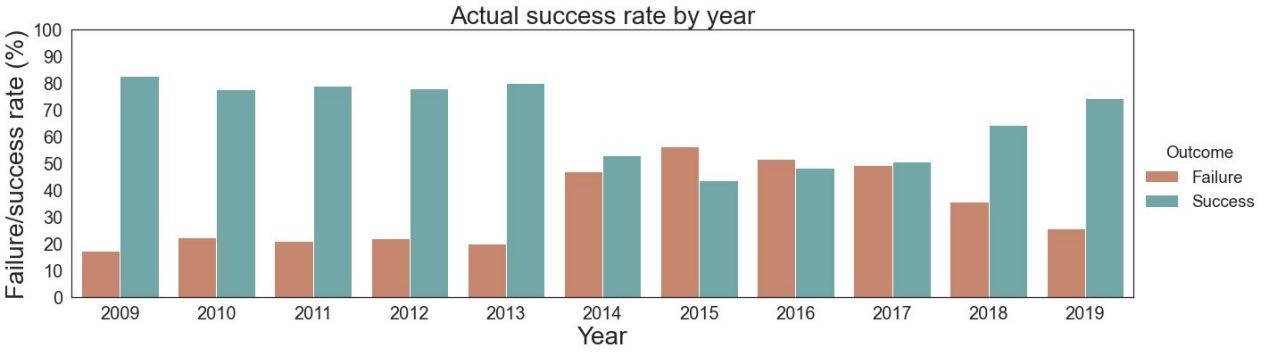

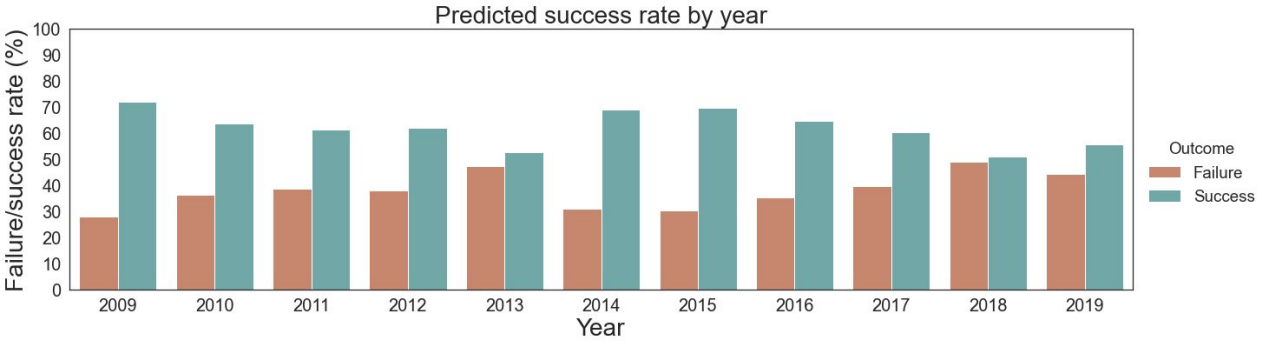

- Actual success rate from 2009 to 2013 is around 80%, which is not captured by the model.
- Reduction in success rate from 2013 to 2014 not captured
- Success rate from 2014 to 2017 is around 50%
- Success rate increases again in 2018 and 2019

Recommendations

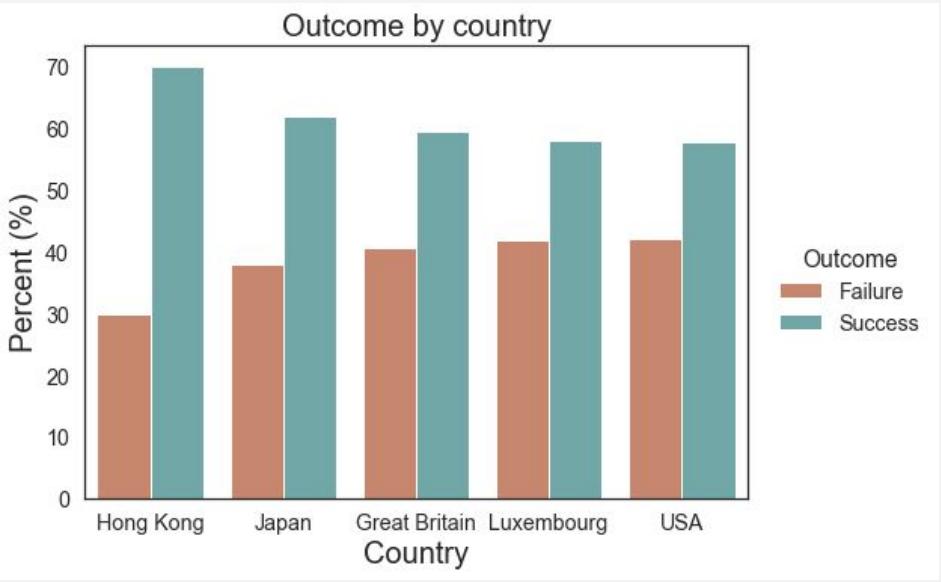

- Projects picked as “staff picks” have a high success rate
- Hong Kong, Japan, Great Britain are top 3 countries with highest success rate

Recommendations

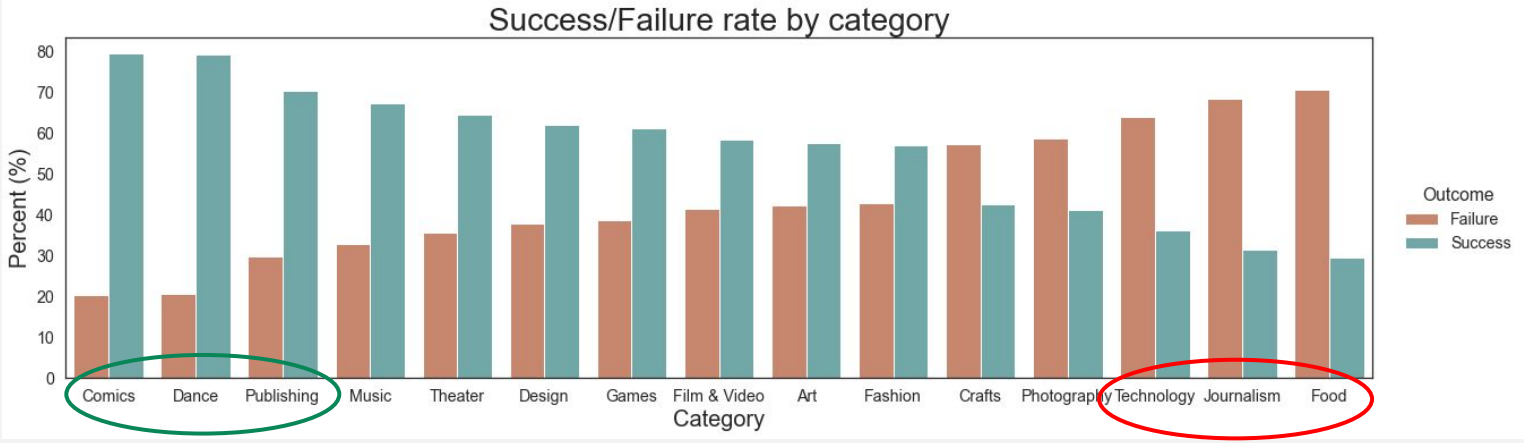

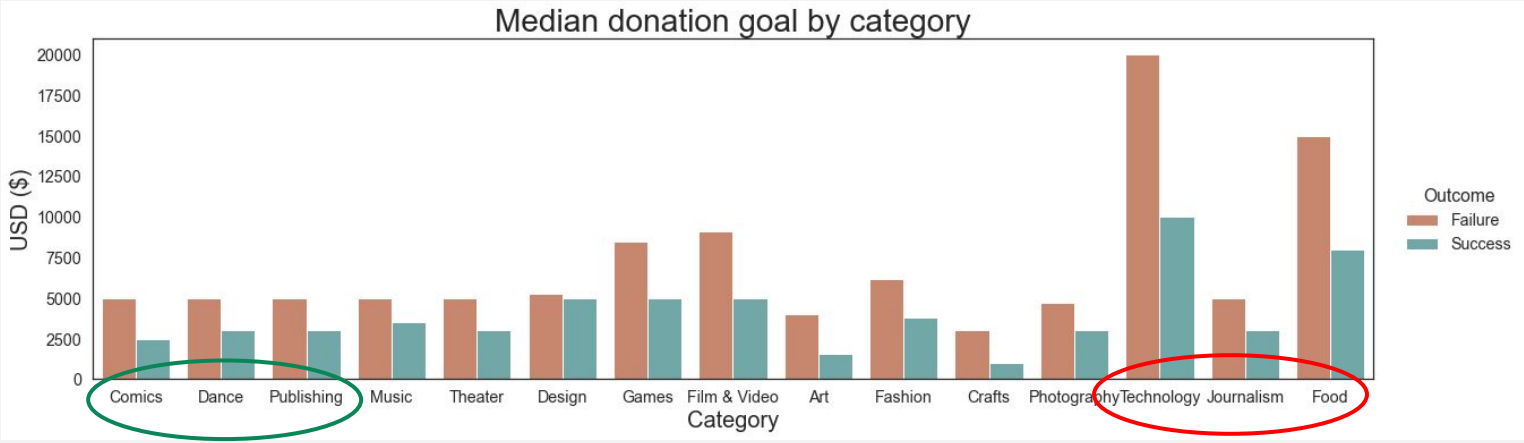

Takeaways and Future Work

- Improvement in accuracy from 55.2% to 72.0% over the baseline model
- Best projects to invest in:
  - Comics, dance or publishing categories
  - Projects featured as “staff pick” when launched
  - Projects from Hong Kong, Japan or Great Britain
- Worst projects to invest in are technology and food and they also have the highest donation goals
- Future work: Explore underlying trends in the data by year or month, subcategory information

Thanks!

Questions?

Erick Cantu Perez
github.com/eaunaicr97

Su Leen Wong
github.com/suleenwong

Project repository:
https://github.com/eaunaicr97/ds-ml-project-kickstarter

Exploratory Data Analysis findings

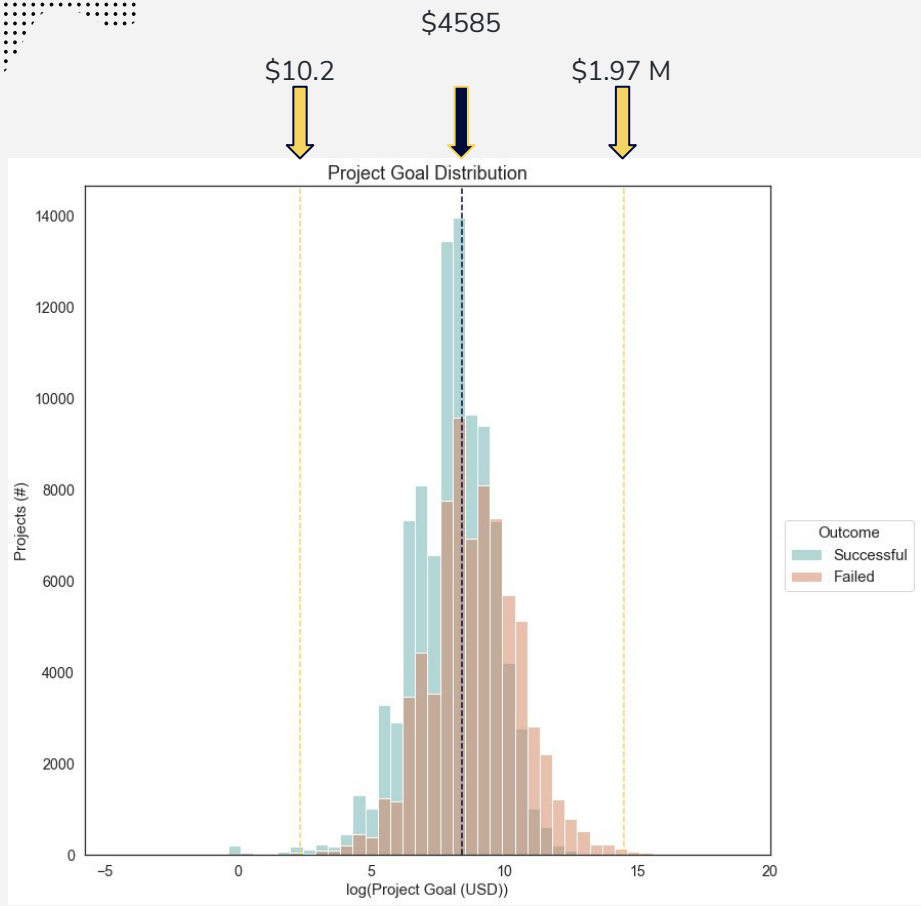

- Failed projects tend to have higher goals
- Project goal log transformed
- 0.523% observations are outside 3.5 std deviations
- 0.002 % successful projects >$1.97
- 0.049 % failed projects < 10.2

Error Analysis

Precision on train data: 0.738
Accuracy on train data : 0.738

Precision on test data: 0.7215
Accuracy on test data : 0.7185

Accuracy on train and test data is quite similar, therefore there is no significant overfitting

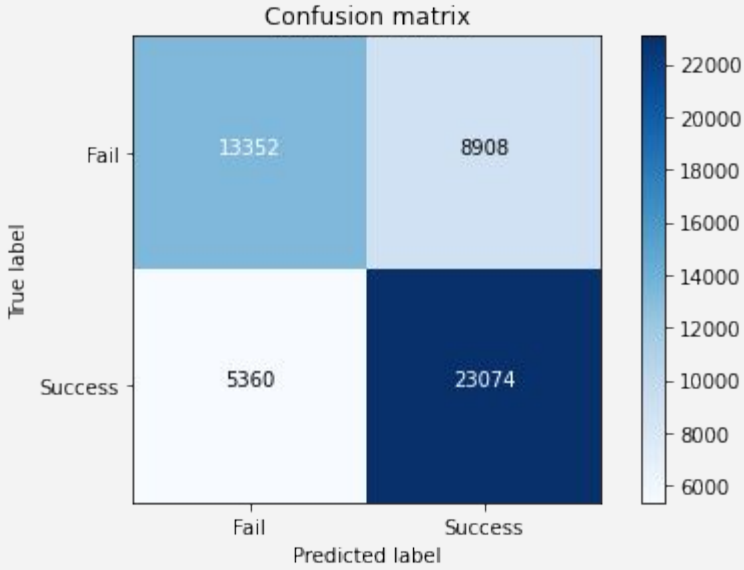

Classification report:

	precision	recall	f1-score	support
0	0.71	0.60	0.65	22260
1	0.72	0.81	0.76	28434
accuracy			0.72	50694
macro avg	0.72	0.71	0.71	50694
weighted avg	0.72	0.72	0.71	50694